
Stage Roadmap: Lexicon, Forward Index & Inverted Index

(Fully Aligned With Your Cleaned Dataset + Industry-Level Design)

Your dataset is already fully cleaned, normalized, tokenized, and refined into `tokens_final`. Now the indexing stage must convert this dataset into fast-searchable structures.

Below is the **exact, optimized, modern roadmap**.

1 STAGE 1 — Lexicon Construction (Vocabulary Dictionary)

Goal: Create a fast, compact mapping:

`token → term_id, document_frequency, pointer_to_postings`

Step 1.1 — Traverse All Documents

Use your `tokens_final` list for each document.

This ensures:

- consistency
 - minimal noise
 - controlled vocabulary (84,787 terms)
 - stable token distribution
-

Step 1.2 — Build the Lexicon (Term Dictionary)

For each distinct token:

Field	Description
<code>term_id</code> (int)	Assigned sequential ID for compact storage
<code>term_string</code>	The token itself
<code>DF (document frequency)</code>	Count of documents containing this term
<code>posting_ptr</code>	Offset pointer to this term's posting list in inverted index

Step 1.3 — Optimize the Lexicon

Apply industry optimizations:

- ✓ Use integer IDs instead of storing strings everywhere
 - ✓ Sort lexicon alphabetically for faster binary search
 - ✓ Optionally compress using FST (light version)
 - ✓ Store DF for BM25 ranking later
-

Step 1.4 — Save Lexicon to Disk

Recommended formats:

- `.lexicon` (binary)
- `.json` for readability during debugging
- `.term_map` for `term_id` → string reverse lookup

This finishes Stage 1.

2 STAGE 2 — Forward Index Construction (Document → Terms)

Goal: Store each document's internal structure for ranking and snippet generation.

Step 2.1 — For Each Document:

Extract:

Field	Description
<code>doc_id</code>	Unique ID (0 → N-1)
<code>term_ids[]</code>	List of term_id for each token
<code>term_freqs</code>	Frequency of each term
<code>positions[]</code>	Optional: token positions for phrase search
<code>length</code>	Document length (used in BM25 normalization)

Step 2.2 — Compress Forward Index

Use:

- Variable-byte encoding
- Delta encoding for positions
- LZ4/Snappy optional compression

Forward index must be lightweight for fast retrieval.

Step 2.3 — Save Forward Index

As segmented files:

- `forward_index_0.bin`
- `forward_index_1.bin`

Segmentation prevents large file rewrite issues.

3 STAGE 3 — Inverted Index Construction (Term → Documents)

This is the heart of the search engine.

Step 3.1 — Build Posting Lists

For each term, create:

Posting = (doc_id, term_frequency, positions[])

Example:

system → [(12,3,[4,10,51]), (47,1,[33]), (490,2,[6,40]), ...]

Step 3.2 — Sort Posting Lists by doc_id

This is required for:

- fast merging
- fast skipping

- fast ranking
-

Step 3.3 — Compress Posting Lists

Industry-standard compression:

- ✓ Delta Encoding (store doc_id differences)
- ✓ Variable Byte Encoding
- ✓ Block-based compression (e.g., 128-doc blocks)
- ✓ Optional: Skip lists

This makes your inverted index **small and extremely fast**.

Step 3.4 — Add Skip Pointers (Highly Recommended)

Add skip pointers every \sqrt{n} entries in the posting list.

Enables:

- fast AND/OR merges
 - fast boolean operator computation
 - fast ranked retrieval
-

Step 3.5 — Save Segmented Inverted Index

Store inverted index in:

`inverted_index_segment_0.bin, segment_1.bin, ...`

Lexicon stores the pointer to each posting list's start.

Stage 4:

1. Purpose First

Stage 4 is all about **preparing the data needed for scoring documents efficiently** at query time.

Think of it as the “fuel” for the ranking engine: you’re not fetching or scoring yet, just computing the constants and tables that will make scoring fast.

- IDF tells you how important a term is across the corpus.
 - Document lengths and average length normalize scores (BM25).
 - Precomputing anything heavy upfront saves milliseconds per query later, which adds up at scale.
-

2. Input / Output Mindset

Inputs:

- Forward index (document → terms, term frequencies, positions)
- Lexicon (DFs for each term)

Outputs:

- `IDF[term_id]` → ready for scoring
- `DL[doc_id]` → document lengths
- `avgDL` → scalar for BM25
- Optional: precomputed BM25 normalization factors like $1 / (k1 * (1-b) + b * DL / avgDL)$

Professional thought: I think in terms of **tables and arrays** that the query engine can load into memory quickly. All these should be numeric, contiguous, and cache-friendly.

3. How I'd Execute

1. Compute IDF:

- Use DF from lexicon: `IDF = log((N - DF + 0.5)/(DF + 0.5) + 1)` (the BM25 version)
- Store in a simple array, indexed by `term_id`

2. Document Length Table:

- Walk forward index, sum term frequencies per document
- Store in an array: `DL[doc_id]`

3. Average Document Length:

- `avgDL = sum(DL) / N`

4. Optional optimizations:

- Precompute `IDF * (k1 + 1)` for BM25 to reduce multiplication during queries
- Store all arrays as **binary files** to minimize loading time

4. Professional mindset

- Stage 4 is **completely read-only** for the indexes; it never modifies the forward/inverted index.
- Everything is **precomputation-heavy**, not query-heavy.
- If done correctly, your query engine just looks up numbers and computes BM25/Tf-IDF without touching large files sequentially.

- This separation of **indexing vs ranking preparation vs querying** is exactly how Lucene, Elasticsearch, or Solr handle it.
-

5. Practical note

I would place **Stage 4** immediately after Stage 3:

Stage 1 → Lexicon
Stage 2 → Forward Index
Stage 3 → Inverted Index
Stage 4 → Ranking Preparation (IDF, DL, avgDL)
Stage 5 → Query Engine

Why: You need DF (lexicon) + document lengths (forward index) to compute ranking metadata.
You can't do ranking prep before Stage 3.

Stage 4 — Query Engine (Search & Retrieval Layer)

Goal: Efficiently answer user queries by ranking and returning the most relevant documents.
This is the **front-facing layer** of your search engine.

Step 4.1 — Query Parsing & Preprocessing

Objective:

Convert raw user queries into a structured form that the search engine can understand.

Tasks:

1. **Tokenization:** Split query into individual words/tokens.

- Example: "best AI books" → ["best", "AI", "books"]

2. Normalization:

- Lowercasing: AI → ai
- Remove punctuation: "books!" → "books"
- Optional: Stemming/Lemmatization (running → run)

3. **Stopword Removal:** Filter out common words that don't affect relevance (the, is, a).

4. **Query Term ID Mapping:** Use your **lexicon** to map tokens → **term_id**.

- Unseen terms → ignored or flagged for expansion.

5. **Phrase Detection (Optional):** Detect quoted phrases ("machine learning") to enforce consecutive positions in results.

Industry Note: Professional engines often also expand queries with synonyms, spell correction, and autocomplete suggestions.

Step 4.2 — Retrieve Candidate Documents

Objective:

Use the **inverted index** to fetch documents containing the query terms.

Tasks:

1. **Lookup Posting Lists:** For each query term, retrieve its posting list from the inverted index.
 - Example: ai → [(doc3, 2, [4, 12]), (doc8, 1, [3])]
2. **Merge Posting Lists:**

- Boolean AND/OR queries: intersect or union posting lists efficiently.
- Use **skip pointers** for fast traversal.

3. Document Scoring:

- Keep a map `doc_id → score` to accumulate relevance scores for each candidate document.

Optimizations:

- Only fetch **top-k candidate documents** if the collection is huge.
 - Skip irrelevant segments using precomputed document frequency or term statistics.
-

Step 4.3 — Ranking & Scoring

Objective:

Rank candidate documents according to relevance using a **ranking function**.

Industry Standard Approaches:

1. Classical Ranking:

- **TF-IDF**: Term Frequency × Inverse Document Frequency.

BM25: State-of-the-art traditional scoring function.

$$\text{Score}(D, Q) = \sum (\text{IDF}(t) * ((\text{TF}(t, D) * (k+1)) / (\text{TF}(t, D) + k*(1-b + b*|D|/\text{avgDL}))))$$

- - $\text{TF}(t, D)$ = frequency of term t in document D
 - $|D|$ = document length

- avgDL = average document length in corpus
- k and b = tunable parameters

2. Optional Modern Enhancements:

- **Learning-to-rank models:** Combine multiple signals (BM25, popularity, CTR).
- **Semantic embeddings:** Map queries & documents into vector space and rank via cosine similarity.

Industry Note: BM25 is the foundation for nearly all traditional search engines like Elasticsearch, Lucene, Solr.

Step 4.4 — Snippet Generation & Highlighting

Objective:

Provide users with a meaningful preview of the document content.

Tasks:

1. **Retrieve token positions** from forward index or inverted index.
2. **Highlight query terms** in context.
3. **Select snippets** that maximize relevance (e.g., first matching paragraph or 2–3 sentences around match).

Industry Standard: Snippets improve user satisfaction and click-through rate.

Step 4.5 — Pagination & Result Formatting

Objective:

Serve results efficiently to users.

Tasks:

- Paginate top-k results.
- Store `doc_id` → `metadata` mapping for title, URL, timestamp.
- Return JSON or HTML depending on application.

Professional Notes:

- Precompute common query results for ultra-fast response (caching).
 - Use **sharding** if your corpus is extremely large, distributing inverted index across servers.
-

Step 4.6 — Query-Time Optimizations

Objective: Minimize latency.

Techniques:

1. **Top-k heap:** Only keep top-k documents while scoring.
 2. **Early termination:** Stop processing low-frequency terms once enough candidates found.
 3. **Parallel Processing:** Query different terms in posting lists concurrently.
 4. **Compression:** Keep inverted index compressed for faster memory access (already in Stage 3).
-

Step 4.7 — Logging & Analytics (Optional but Professional)

- Track queries, click-through rate, and term popularity.

- Helps improve ranking algorithms and autocomplete suggestions.
-

Stage 4 Flow — Visualized (Conceptual)

USER QUERY

|

v

[Preprocessing: tokenize, normalize, stopwords]

|

v

[Map tokens → term_id using Lexicon]

|

v

[Retrieve Posting Lists from Inverted Index]

|

v

[Merge Posting Lists → Candidate Docs]

|

v

[Score Documents (BM25 / TF-IDF / Learning-to-rank)]

|

v

[Generate Snippets / Highlight Query Terms]

|

v

[Sort & Return Top-k Results to User]

Professional Industry Notes

1. Separation of Concerns:

- Stage 1–3 = offline indexing.
- Stage 4 = online query engine.
- Professional engines separate these layers for scalability.

2. Performance Considerations:

- Query response < 200ms for web-scale engines.
- Use memory-efficient data structures (compressed posting lists, skip pointers).

3. Optional Industry Enhancements:

- Caching popular queries.
 - Distributed query execution across shards.
 - Semantic search (vector embeddings + ANN search).
-

5 STAGE 5 — Validation

Before moving ahead:

- ✓ **Check lexicon size matches 84,787**
- ✓ **Ensure total postings = sum of all term frequencies**
- ✓ **Random search queries to test:**

- system
 - iphone
 - education
 - siberia
 - quantum
 - mango
-

6 Final Output of this Stage

After finishing, you will have these files:

File	Purpose
lexicon.bin	Term dictionary, DF, posting pointers
term_map.json	term_id → string reverse map
forward_index.bin	Document → term_ids
doc_length.bin	Document lengths
inverted_index.bi n	Posting lists for all terms
stats.json	Global stats (N, avg_doc_len, vocabulary size)

These form the **search engine's backbone**.

7 Why This Roadmap Is Industry-Perfect

Because it follows real world search engine design principles:

- ✓ Lucene architecture
- ✓ Elasticsearch indexing
- ✓ Vespa streaming indexes
- ✓ Bing/Yandex block-max index design
- ✓ Google's dictionary + posting segmentation approach

Your dataset (cleaned, normalized, 84k-term lexicon) fits perfectly into this architecture.

```
search_engine_project/
├── data/
│   └── sample_tokens_final_clean.txt
├── src/
│   ├── main.cpp
│   ├── stage1_lexicon.cpp
│   ├── stage2_forward_index.cpp
│   └── stage3_inverted_index.cpp
└── include/
    ├── stage1_lexicon.h
    ├── stage2_forward_index.h
    └── stage3_inverted_index.h
└── output/
    ├── lexicon/
    ├── forward_index/
    └── inverted_index/
└── build/
    └── search_engine.exe
```