

Stage 4 — Ranking Precompute (Search Engine Project)

Date: 2025-11-25

Project: Custom Search Engine

Stage: 4 / Phase A — Productionize Single-Node

Overview

Stage 4 focuses on **precomputing ranking-related statistics** to make the query engine efficient and scalable. Specifically, it calculates:

1. **Document lengths (`d1`)** — number of terms per document.
2. **Average document length (`avgd1`)** — used in BM25 normalization.
3. **Inverse Document Frequency (`idf`)** — measures term importance across the corpus.

These outputs are **binary-serialized** for fast loading in the next stage (Stage 5 — Query Engine).

Inputs

- **Dataset:** `data/sample_tokens_final_clean.txt` (cleaned tokenized documents)
 - **Lexicon:** Built in Stage 1 (`stage1_lexicon`) → maps terms to `term_id` and stores document frequency.
 - **Forward Index:** Built in Stage 2 (`stage2_forward_index`) → stores per-document term lists and frequencies.
-

Process

1 Loading Lexicon and Forward Index

- **Lexicon:** Provides `term_id` mapping for all tokens in the corpus.
- **Forward Index:** Provides per-document term frequencies and lengths.

2 Compute Document Lengths (`dl`)

For each document in the forward index:

```
dl[doc_id] = total number of terms in document
```

This is used for BM25 normalization (`b` parameter).

3 Compute Average Document Length (`avgdl`)

$$\text{avgdl} = \sum_{i=0}^{N-1} \text{dl}[i] / N$$

Where `N` is the total number of documents.

Result in project: `avgdl ≈ 318.471` for 5000 documents.

4 Compute Inverse Document Frequency (`idf`)

BM25-style IDF is calculated as:

$$\text{idf}(t) = \log N - \text{df}(t) + 0.5 \text{idf}(t) = \log \frac{N - \text{df}(t) + 0.5}{\text{df}(t) + 0.5}$$

Where:

- `df(t)` = number of documents containing term `t`
- `N` = total documents

This captures **term importance**: rarer terms get higher IDF.

Implementation Highlights

- Written in **C++**, integrated with existing project structure:
 - `include/stage4_ranking.h`
 - `src/stage4_ranking.cpp`
 - Modular design: `Stage4Ranking` class accepts `ForwardIndex` and `Lexicon` objects.
 - Binary serialization for **fast Stage 5 loading**:
 - `d1.bin` → vector of document lengths
 - `avgdl.bin` → single double value
 - `idf.bin` → map of `term_id` → IDF
 - Files saved to: `output/ranking/`
-

Execution

In `main.cpp`, Stage 4 is executed after Stage 3:

```
Stage4Ranking stage4(fwd, lex);
stage4.compute();
stage4.save_binaries(output_folder + "\\ranking");
```

Console Output:

```
Stage 4: Ranking precompute done.
Total docs: 5000, avgdl: 318.471, total terms: 46763
Stage 4: Binaries saved to output\ranking
```

- Processed **5000 documents**

- 46763 unique terms
 - Computed avgdl = 318.471
-

Deliverables

- `output/ranking/dl.bin` — document lengths
 - `output/ranking/avgdl.bin` — average document length
 - `output/ranking/idf.bin` — term IDF
 - Fully ready for Stage 5 Query Engine
-

Challenges & Solutions

Challenge	Solution
Handling large datasets efficiently	Used binary files for fast read/write
Mapping term → term_id consistently	Used <code>Lexicon</code> from Stage 1
Computing document frequency efficiently	Used per-document unique term sets to avoid double counting
Integrating with existing project structure	Designed <code>Stage4Ranking</code> class compatible with <code>ForwardIndex</code> and <code>Lexicon</code>

Next Steps

- Stage 5 — Query Engine
 - Load `dl.bin`, `avgdl.bin`, and `idf.bin`

- Implement BM25 scoring
- Top-K retrieval
- Snippet generation
- Serve queries via a REST API

Stage 4 ensures that **all ranking statistics are precomputed**, enabling the query engine to perform **fast, accurate searches** over 5000+ documents with minimal runtime computation.